

UNIVERSITY OF TARTU
Institute of Computer Science
Computer Science Curriculum

Tiit Oja

Optimizing JVM profiling performance for Honest Profiler

Bachelor's Thesis (9 ECTS)

Supervisor: Vootele Rõtov, MSc

Supervisor: Vesal Vojdani, PhD

Tartu 2018

There's stuff to do!

Mare Koit suggested to do center alignment for text paragraph	3
The first part of the sentence does not connect very well with the second. Maybe it is because of “results” and “tools” being generic terms whose meaning one has to parse with context in mind. Reads have only one paragraph worth of context by time they read this	6
Bit thorn here, we are guaranting that increased samples increases accuracy, but this might be overselling it. At same time, where else to oversell if not in the introduction? Using some sort of conditonal there would make it much weaker. Maybe ask for a second opinion from someone on this.	6
Verify that this code snippet is completely rendered on the same page in the final version	10
Nested loop in the example is confusing, simplify and reproduce	10
Can just add a comment	10
Ask Vootele what comment he means here?	10
Check phrasing of the following paragraph with supervisors	11
Could use an example, some source code before and after instrumentation, on Java code level (added stopwatch start and stop or something) . . .	12
Add code sample	12
Cite signal safety man page	13
I would move this second, it is way more common than AsyncGetCallTrace profilers and seems to make sense to introduce it first? But maybe just personal preference.	13
Verify that this listing is entirely on the same page before shipping final version	16
Before shipping final, verify that the pictures are positioned well	18
Is defining test machine specs necessary?	18
Is it better to move benchmark profiling result visualizations to appendices?	18
Signal handler safety worth mentioning?	23
sleeping logic to reduce overhead worth mentioning?	24
Should present all cases or most common ones would suffice?	26
and in which use-cases it outshines the others – sampling bias etc	27
other	27
First sentence in past, second in present. Also, style question, in the thesis or thesis?	27
there was no fundamental reason why it was not possible?	27
Discussed in Slack. I think currently, the two other solutions are undersold. .	27
Add date	30

Optimizing JVM profiling performance for Honest Profiler

Abstract: Honest Profiler is a profiling tool which extracts performance information from applications running on the Java Virtual Machine. This information helps to locate the performance bottlenecks in the observable application. This thesis aims to provide solutions to increase the output of useful information by Honest Profiler. Achieving this would increase the efficiency and accuracy of the profiling results collected by Honest Profiler. Thesis will cover the basics of sampling profiling, the architecture of Honest Profiler and measures the performance of Honest Profiler's data collection logic. As a main result of this thesis, three different solutions for increasing the profiler information output are presented and their performance and the extracted information amount is evaluated by a benchmark test.

Mare Koit suggested to do center alignment for text paragraph

Keywords: Profiling, optimization, Honest Profiler

CERCS: P170, Computer science, numerical analysis, systems, control

JVM profileerimise jõudluse optimeerimine Honest Profileri baasil

Lühikokkuvõte: Honest Profiler on tööriist, mis võimaldab mõõta Java virtuaalmasina peal jooksvate rakenduste jõudlust. Tööriista poolt kogutud informatsioon annab ülevaate rakenduse jõudlusest, mille baasil on võimalik optimeerida vaadeldava rakenduse jõudlust. Käesoleva töö eesmärk on luua lahendusi, mis suurendaksid Honest Profileri tööriista poolt kogutud informatsiooni hulka. Suurem kogutud informatsiooni hulk muudab jõudluse mõõtmise efektiivsemaks ning saadud tulemused täpsemaks. Töö kirjeldab profiilide kogumise ning Honest Profileri arhitektuuri põhitõdesid ning mõõdab Honest Profileri informatsiooni kogumise loogika jõudlust. Töö põhitulem on kolm erinevat lähenemist, mis suurendavad kogutud informatsiooni hulka. Kirjeldatud lahenduste jõudlus ning täpsus verifitseeritakse jõudlustesti abil.

Võtmesõnad: Jõudluse mõõtmine, optimeerimine, Honest Profiler

CERCS: P170, Arvutiteadus, arvutusmeetodid, süsteemid, juhtimine (automaatjuhtimisteooria)

Contents

1	Introduction	6
2	Sampling profiling methodologies	7
2.1	Importance of samples' quantity	8
2.2	Sampling profiling challenges	9
2.2.1	Periodicity bias	9
2.2.2	Safepoint bias	10
2.2.3	Observer effect	11
2.3	Sampling profiling compared to alternatives	12
2.4	Sampling profiling possible implementations for the JVM	12
2.4.1	jstack	12
2.4.2	AsyncGetCallTrace profilers	13
2.4.3	GetAllStackTraces profilers	13
2.4.4	Native profilers	14
3	Honest Profiler	15
3.1	Architecture	15
3.2	Current limitations	16
3.3	AsyncGetCallTrace performance analysis	16
3.4	Implications of the performance analysis	17
4	Profiler optimization	18
4.1	Benchmark	18
4.2	Optimization implementations	19
4.2.1	Interval timer based on real time	19
4.2.2	Increasing the kernel clock frequency	21
4.2.3	External utility for timekeeping	22
5	Future research	26
6	Conclusion	27
	Appendices	29
A	Source code	29
B	Benchmark test code	29

1 Introduction

Profiling is an activity which aims to identify performance issues in an observable program. This task often relies on specific tools called profilers for extracting information from the program's execution. The output of the profiler helps to identify and locate the methods that have the largest effect on the program's execution time. Such methods are usually worth investigating as they affect the application's performance the most. [11]

As the profiling results rely on the efficiency and accuracy of the tools used,

The first part of the sentence does not connect very well with the second. Maybe it is because of “results” and “tools” being generic terms whose meaning one has to parse with context in mind. Reads have only one paragraph worth of context by time they read this

it is important to have sufficient and actionable information from the profiler to construct accurate profiles of the application's performance. This thesis investigates a specific profiler implementation named Honest Profiler which is an open source sampling profiling tool which can be used to evaluate the performance of Java applications. The goal of this thesis is to increase the amount of information extracted from an observable Java application by Honest Profiler. Doing so increases the accuracy and reliability of the results.

Bit thorn here, we are guaranting that increased samples increases accuracy, but this might be overselling it. At same time, where else to oversell if not in the introduction? Using some sort of conditonal there would make it much weaker. Maybe ask for a second opinion from someone on this.

This thesis covers the basics of sampling profiling methods and the problems these methods have in the context of Honest Profiler. It explains the architecture of Honest Profiler and measures the performance of its profiling logic. The main result of this thesis is providing means to increase the amount of information that Honest Profiler can extract from the observable application. The suggested solutions are then tested on a benchmark test to evaluate their performance and amount of useful information extracted.

2 Sampling profiling methodologies

Sampling profilers are a type of profiles which gather call traces from the observable program at varying intervals. A sample in the form of a call trace is a representation of a single thread's state at a particular moment in time. A simple call trace of a thread can be observed in listing 1.

```
"main" #1 prio=5 os_prio=0 tid=0x00007feccc224000 nid=0x1f2c
runnable [0x00007fec5660000]
  java.lang.Thread.State: RUNNABLE
    at ee.ut.SimpleBenchmark.methodA(SimpleBenchmark.java:28)
    at ee.ut.SimpleBenchmark.doWork(SimpleBenchmark.java:22)
    at ee.ut.SimpleBenchmark.main(SimpleBenchmark.java:11)
```

Listing 1: Call trace of a thread

Gathered samples can then be aggregated and grouped based on their occurrence. Distribution of the gathered samples highlights the hotspots in the observable program. Higher frequency of a call trace suggests that more program's execution time was spent in that particular state.

Figure 1 helps to visualize the performance implications of the frequency of the gathered samples. The visualization was created by a tool called Flamegraphs which provides means to generate a comprehensive and intuitive visualization of the gathered call trace samples. [10] Visualization tools are necessary to produce meaningful representations of the collected data as unprocessed call trace samples *per se* are not informative without the context of other samples' frequency.




Figure 1: Visualization of sampling profiling output

The samples of the visualization in figure 1 were gathered from a simple program that executed two identical methods, `methodA` and `methodB`, alternately. During its execution 481 usable samples were collected. 198 of those samples contained `methodA` in its top call frame and 194 contained `methodB` in its top frame.

Although sampling profiling does not provide precise metrics for each method's execution time, it can provide a general overview to visualize the time spent in the profileable application's context. Such information is often sufficient to make actionable optimizations in the observable application.

2.1 Importance of samples' quantity

For sampling profilers, the result is a statistical approximation of the program's performance. Thus, having more samples will provide a more accurate approximation.

Figure 2 on page 8 illustrates the problem which is caused by the low sampling interval. Suppose that the figure 2 resembles a program's execution in which the horizontal axis represents the current program's call trace state in that particular moment in time. It can be observed that the amount of time spent in method Y is significantly larger than the amount of time that was spent in methods X and Z. Suppose that the sampling profiler obtains a sample at each dotted vertical line. Such profiling results would represent each method X, Y and Z with a single sample. This implies that all the methods spent roughly the same amount of time during the execution of the program. The result is inaccurate as the visualization clearly shows that method Y spent roughly 4 times more time than it was spent for executing methods X and Z.




Figure 2: Scenario with low sampling interval

Increasing the sampling interval would improve the situation as demonstrated on figure 3 on page 9. Upon profiling with a four times higher sampling interval, it becomes apparent that method Y call trace samples have been proportionally represented in the total call trace samples. The call trace samples containing the method Y in its top frame are colored red on the figure 3.




Figure 3: Scenario with sufficient sampling interval

2.2 Sampling profiling challenges

There are multiple problems that the implementations of sampling profilers could experience. These problems could potentially affect the reliability of the samples provided by the sampling profiler.

2.2.1 Periodicity bias

This problem occurs when the sampling interval catches on to some program's routine which executions match the sampling interval. Figure 4 illustrates the issue behind the bias. Suppose that the observable program runs **Method X** and **Method Y** alternately for some constant time period. If the dotted lines are the marks for the call trace samples taken during profiling, the results would be skewed as not a single sample represents **Method Y** in the results.




Figure 4: Illustration of periodicity bias

Possible way to tackle this problem would be to randomize the sampling interval by having a random number of time units offsetting the sampling interval. Increasing sample gathering frequency will also help to alleviate this problem.

2.2.2 Safepoint bias

Sampling profiling assumes that the obtained samples are acquired uniformly and randomly as bias in the samples could potentially yield inaccurate results. In light of this, when gathering samples from a program running on the Java Virtual Machine, one must take *safepoints* into consideration.

Safepoints in the Java Virtual Machine are defined as points during the program's execution during which all of the threads are in a consistent and well known state. Safepoints are necessary for various Java Virtual Machine's operations such as the garbage collection, method deoptimization and class redefinition.[3] It has been shown that most of the existing Java profilers require the Java Virtual Machine to be stopped on a safepoint in order to obtain the call trace sample from the profileable thread.[15] However, this mechanism also casts a shadow on the profiling results as this implies that the samples are not acquired randomly but rather require the Java Virtual Machine to be in a specific state. [11]

The execution of the code sample in listing 2 illustrates this issue well.

Verify that this code snippet is completely rendered on the same page in the final version

Nested loop in the example is confusing, simplify and reproduce

Can just add a comment

Ask Vootele what comment he means here?

```
int k = 0;
for (int i = 0; i < Integer.MAX_VALUE; i++) {
    for (int j = 0; j < 2; j++) {
        k++;
        if ((k % 2) == 1) k++;
    }
}
```

Listing 2: Counted loops do not contain safepoints

As various Java Virtual Machine routines are nondeterministic, it is impossible to predict when does the Java Virtual Machine signal its threads to be stopped on

a safepoint. However, if such request should occur in the middle of the counted loop's execution in listing 2, the thread executing the loop's instructions will not stop until the loop has finished. Thus, if an ordinary sampling profiler signals the thread for a call trace sample, it is delayed until the thread executing the loop's instructions finishes its task. [14]

To measure the actual time that is spent waiting for the Java Virtual Machine to stop all its threads on a safepoint, `-XX:+PrintGCApplicationStoppedTime` flag can be used. Appending this flag to the Java Virtual Machine arguments, it enables outputting the time it takes for the Java Virtual Machine to stop all threads in order to execute a subroutine and the time that this operation took in total. Due to nondeterminicity, the sample code in listing 2 was executed 100 times. The worst case among 100 program runtimes recorded a Java Virtual Machine operation that took 7.517 seconds and 99.999% of that time was spent on reaching a safepoint. It is worth mentioning that this operation took 65% of the whole program's execution time.

Although the provided example is rather artificial and such occurrences in real life programs are rare on such large scale, it clearly points out the bias that the safepoints potentially introduce. Such bias during obtaining the call trace clearly contradicts the randomness prerequisite for sampling profiling. Despite the presence of safepoint bias, such profiler can still produce actionable profiling results but lessen the accuracy when high detail granularity of the results is important.

2.2.3 Observer effect

Attaching a profiler to the observable application does alter the way how the application would normally execute. The *observer effect* describes how a presence of a profiler can change the profiling outcome.

Firstly, a profiler does introduce some performance overhead due to its nature. Profilers do need to execute additional instructions for gathering information from the observable application. To have meaningful profiling results, this overhead must not affect the application's performance by a large margin.

Check phrasing of the following paragraph with supervisors

The presence of a profiler could also change the way how the Java Virtual Machine does its optimizations. Various information extracting functions might change the way how the bytecode on the Java Virtual Machine is interpreted. [11] Thus, resulting in a different performance profile when compared to an application's runtime in which the profiler is not attached.

2.3 Sampling profiling compared to alternatives

An alternative approach for gathering performance information is using tracing profilers. Tracing profilers rely on observable application's bytecode instrumentation. Instrumentation is the act of adding additional instructions to the observable application to gather relevant profiling information. Tracing profilers instrument the observable application by wrapping method calls with instructions to measure method execution time. Result is a collection of call stacks with included precise method timings. [16]

Could use an example, some source code before and after instrumentation, on Java code level (added stopwatch start and stop or something)

Add code sample

Using tracing profilers usually requires the user or the profiler to know the application elements that are worth profiling as instrumenting all the methods in the observable application is too expensive performance wise.

Downside of tracing profiling is its performance and its susceptibility to observer effect. Tracing profilers rely on bytecode instrumentation which directly affects the way how the Java Virtual Machine performs its optimizations. Some optimizations might not be possible due to the added bytecode instructions added by the tracing profiler. Additionally, depending on the level of information detail and the amount of instrumentation performed, requiring higher information granularity could potentially introduce impactful overhead to the application's performance which could result in inaccurate profiling results.

When comparing sampling profiling to tracing profiling, it excels with its performance. Sampling profilers can potentially obtain thousands of samples each second with negligible performance overhead.

2.4 Sampling profiling possible implementations for the JVM

2.4.1 `jstack`

`jstack` is an utility included in the Oracle's JDK and OpenJDK by default. This utility can be used to print out stack traces of all threads of a Java process.[4] The simplest implementation to show the concept of a sampling profiler would be to periodically call the `jstack` utility to obtain the snapshot of all threads.

These snapshots can be persisted and processed to produce a visualization of the application's performance. This approach has a relatively high overhead when compared to other sampling profiling methods and is not a viable method for profiling real world applications but rather illustrates the concept of sampling profiling.

2.4.2 `AsyncGetCallTrace` profilers

Cite signal safety man page

These kind of profilers make use of an undocumented JVM method `AsyncGetCallTrace` [7] which enables obtaining call traces from a thread without the safepoint bias. The asynchronous nature of this method enables the profiler to safely call this method in a signal handler which is the key architectural element for these kind of profilers.

Due to safepoint bias free sample obtaining method, profiling samples tend to be more accurate since the JVM does not have to stop on a safepoint in order to obtain the call trace of the running thread. This method is a part of the JVM Tooling Interface (JVMTI) Application Programming Interface (API).

Notable examples of such profilers which utilize `AsyncGetCallTrace` method are Honest Profiler and Java Mission Control.

2.4.3 `GetAllStackTraces` profilers

I would move this second, it is way more common than `AsyncGetCallTrace` profilers and seems to make sense to introduce it first? But maybe just personal preference.

`GetAllStackTraces` is a method in the official JVMTI API which can be used to extract the stack traces of all the threads currently executing on the JVM[5]. Downside for this method is the fact that profilers utilizing this method are safepoint biased. For this reason, the performance is negatively affected due to frequent pauses for reaching safepoints.

Due to safepoint induced performance problems, this method acquired significantly less samples when compared to profilers based on `AsyncGetCallTrace`. Additionally, profiling results might be biased due to the fact that the stack traces can only be obtained when the JVM has reached a safepoint.

`GetAllStackTraces` is used in most commercial profilers such as JVisualVM, YourKit and JProfiler. [9] [15]

2.4.4 Native profilers

Native profilers are tools utilizing the low level abstractions of the operating system to profile the native binaries for the operating system. These profilers are superior to the previously described methods performance wise but do fall short on the usefulness of the gathered information. The main downside of native profilers is the inability obtain sufficient information about the Java level stack frames executed inside the Java Virtual Machine. Unlike other sampling techniques introduced, native profilers are not much use to debug a JVM where most of the code has not been yet optimized

I would not bring in more here, let's leave it out of scope. Maybe add that “unlike other sampling techniques introduced, native profilers are not much use to debug a JVM where most of the code has not been yet optimized” or something similar, to give a reason why we leave those out.

3 Honest Profiler

Honest Profiler [2] is a software written in Java and C++ which aims to gather honest and accurate samples from a running Java application to provide an overview of the application’s performance. Honest Profiler uses the `AsyncGetCallTrace` method for obtaining the samples. Moreover, Honest Profiler is licensed under an open source license which encourages the community to further investigate and develop this software solution. The fact that Honest Profiler is an open source software and utilizes a safepoint bias free method for obtaining the samples are the main reasons why this particular software is chosen for improvements and optimizations in this thesis.

3.1 Architecture

Honest Profiler consists of two larger components: a JVMTI native agent and a Java facade for result visualization, analysis and transformation. The profiling is almost completely done by the agent whereas the Java facade component aims to provide tools to ease the use of the profiler.

The JVMTI agent is a dynamic library which serves as a client of the JVM Tooling Interface. [6] JVMTI API provides these agents means to investigate and control the applications running on the Java Virtual Machine. Since the agents are written in native languages such as C and C++, they can also utilize various operating system’s tools.

Honest Profiler relies on UNIX operating system signals and timers for obtaining samples. It sets up a signal handler to handle `SIGPROF` signals sent to the Java process that the agent is attached to and an interval timer which sends the `SIGPROF` signal to that same process periodically.

Upon receiving a signal, an *arbitrary* thread currently executing on the CPU will handle the signal. [12] The signal handler proceeds to call `AsyncGetCallTrace` method from the JVMTI API to obtain the call trace of the current thread which is handling the signal. The obtained call trace is then persisted.

It is worth mentioning that the thread handling the `SIGPROF` signal can be any thread of the Java process being profiled. Even the threads responsible for garbage collection and *just-in-time* compilation. In such cases, the samples are discarded as such samples do not provide informative results.

3.2 Current limitations

Current implementation of Honest Profiler makes use of the standard interval timer from GNU time library and UNIX operating system signals to periodically call the `AsyncGetCallTrace` method. This implies that the lowest possible `AsyncGetCallTrace` method call interval is the time between two ticks of the system timer interrupt (*jiffy*). This duration depends on the clock interrupt frequency of the hardware platform being used. Default clock frequency for Linux based operating systems (since kernel version 2.6.13) is 250 Hz which results in $\frac{1}{250} = 0.004$ seconds for the duration of a jiffy. [8] In such configuration, minimum sampling interval would be 4 milliseconds.

3.3 AsyncGetCallTrace performance analysis

Prior to increasing the sampling frequency, it is important to understand to what are the performance limitations that calling the `AsyncGetCallTrace` would imply. It is necessary to understand how long does a call to the `AsyncGetCallTrace` method take to ensure that increasing the sampling frequency will not cause excessive performance overhead to the application being profiled. In order to measure the execution time, a custom OpenJDK build with CPU time based time measurement logic was added to the `AsyncGetCallTrace` method. The added timing functionality can be observed in listing 3.

Verify that this listing is entirely on the same page before shipping final version

```
void AsyncGetCallTrace(ASGCT_CallTrace *trace, jint depth, void* ucontext) {
    struct timespec start, stop;
    clock_gettime(CLOCK_THREAD_CPUTIME_ID, &start);

    // AsyncGetCallTrace method code

    clock_gettime(CLOCK_THREAD_CPUTIME_ID, &stop);
    long accum = stop.tv_nsec - start.tv_nsec;
    printf("%li", accum);
}
```

Listing 3: CPU time based measurement in `AsyncGetCallTrace` method

Testing shows that a single `AsyncGetCallTrace` method call takes 4.6 microseconds on average.

3.4 Implications of the performance analysis

Previously described limitations and `AsyncGetCallTrace` method performance measurements do not expose apparent problems or reasons on why `AsyncGetCallTrace` method can not be called more frequently during profiling. As the `AsyncGetCallTrace` method execution time is significantly smaller when compared to the default lowest achievable sampling interval, it can be safely assumed that having smaller profiling interval will not affect the observable application's performance to a great extent. Thus, achieving higher sampling frequency is a goal worth pursuing.

4 Profiler optimization

This sections describes three different approaches for increasing the sampling frequency for Honest Profiler. To verify the performance overhead of the solution and its accuracy, a benchmark test was executed on each solution.

4.1 Benchmark

Benchmark test reads 10000000 random integers from an input file and then performs quicksort on the copy of the obtained list of integer 10 times. Benchmark test code is included in appendix B. All tests are performed on a machine running Intel i5-5200U processor with 8 GB of RAM.

Baseline of the test using the default implementation of the Honest Profiler finished the test in 74.72 seconds on average. During profiling, 19025 samples were obtained on average from which 15130 (81%) contained frames of interest. Profiling results are visualized on the figure 5 on page 19.

Before shipping final, verify that the pictures are positioned well

Is defining test machine specs necessary?

Is it better to move benchmark profiling result visualizations to appendices?




Figure 5: Visualization of baseline benchmark profiling results

4.2 Optimization implementations

This thesis covers three methods to obtain more samples from the observable application than it would be for the default implementation of the Honest Profiler. Following sections will describe their implementations more in-depth.

4.2.1 Interval timer based on real time

Existing implementation of Honest Profiler uses CPU time based interval timer from GNU time library. Alternative approach would be to use an interval timer which counts time in real time instead. The very same interval timer from GNU time library can be configured in such way that it counts down in real time by initializing it with `ITIMER_REAL` flag in the manner as shown in listing 4.

```
#include <sys/time.h>
```

```
struct itimerval timer;  
setitimer(ITIMER_REAL, &timer, 0)
```

Listing 4: Timer initialization based on real time

Upon timer expiration, a **SIGALRM** signal is sent to the process.[1]

Using real time based timer means that its interval is not limited by the system's clock frequency as it would be for CPU time based timers. Due to this fact, the timer can even be configured to send the **SIGALRM** signal every 1 microsecond.

Although this approach obtains considerably larger amount of samples due to higher profiling frequency, it also introduces a potential bias that could affect the profiling results. It might be possible that between the two consecutive timer expiration signals, the observable application has not received any processor time to perform its instructions. This would imply that the same sample in the exact same state is persisted multiple times. Similar issue occurs the other way around when the application's instructions are executed on multiple processor cores. There is a clear distinction in cases in which 4 processor cores are fully utilized executing the application's instructions and cases in which a single processor core is used. The former scenario would get roughly four times as much work done when compared to the latter scenario. However, the real time based timer will not take these workload differences into account because it measures time based on the real world clock.

Average benchmark runtime using this method finished in 82.36 seconds and obtained up to 12.8 million samples during its runtime from which only 73982 (0.58%) contained frames of interest. Whereas the 10% increase in execution times is acceptable for gaining significantly larger amount of samples, it is unknown why only 0.58% of the gathered samples contained frames of interest. This phenomenon is also further covered in the future research section 5. Visualization of the profiling results utilizing this method can be observed in figure 6 on page 21.




Figure 6: Visualization of real time based solution profiling results

4.2.2 Increasing the kernel clock frequency

As CPU time based interval timer from GNU timer library is directly dependent on the kernel's internal timer frequency, one approach to gather more samples to build a custom kernel with increased kernel clock frequency. By default this frequency is set to 250 Hz.[13] This value can be increased to achieve higher sampling frequency.

This can be accomplished by creating a custom configuration in the `Kconfig.hz` with a desired frequency value and building the kernel using the created configuration. An easier way to obtain a kernel with higher kernel timer frequency would be to install a `linux-lowlatency` kernel package which has set the kernel timer frequency to 1000 Hz.

Using the 1000 Hz kernel increases the sampling frequency as the CPU time based interval timer from GNU time library can now send an interrupt signal every $\frac{1}{1000} = 0.001$ seconds. Experiments with such kernel has shown four times increase in the amount of samples gathered without causing significant performance overhead to the observable application.

Benchmark runs on such configuration had average runtime of 77.33 seconds during which 70525 samples were collected. 50937 of the samples (72%) consisted of meaningful samples. When comparing to the baseline benchmark results, this solution resulted in 336% increase in useful samples with negligible performance

overhead. Visualization from a runtime using such profiling configuration is shown on figure 7 on page 22.




Figure 7: Visualization of the results from profiling with increased kernel timer frequency

4.2.3 External utility for timekeeping

This approach moves the time measurement logic to an external server process which uses UNIX operating system signals and shared memory to communicate with Honest Profiler and the profileable Java process. The solution is based on the fact that it is possible for a process to obtain the CPU time that the process has used without the limitations imposed by the kernel's internal timer frequency. The CPU time spent by a process can be obtained by calling the `clock_gettime` function from the GNU time library.

This solution has two benefits over the default implementation of Honest Profiler:

1. Extracting the interval timing logic from the Honest Profiler native agent reduces performance overhead caused to the observable Java process
2. Enables calling `AsyncGetCallTrace` method more often than default implementation

The external server process starts by creating an UNIX domain socket and waits for a connection. Upon a connection from the observable Java process with Honest Profiler's native agent attached, it obtains its process identifier (*pid*) for signal sending purposes. The server process then allocates a memory section which is shared with the profileable Java process. Server will then use two signals to request action from the Honest Profiler that is attached to the Java process:

Signal handler safety worth mentioning?

- `SIGALRM` - write the used CPU time used by the Java process to the shared memory
- `SIGPROF` - call the `AsyncGetCallTrace` method to obtain a sample from the thread handling the signal

Shared memory for interprocess communication is used due to its superior performance over other communication approaches such as file I/O (input/output) or UNIX domain sockets.

Figure 8 on page 24 summarizes the architecture of this implementation.

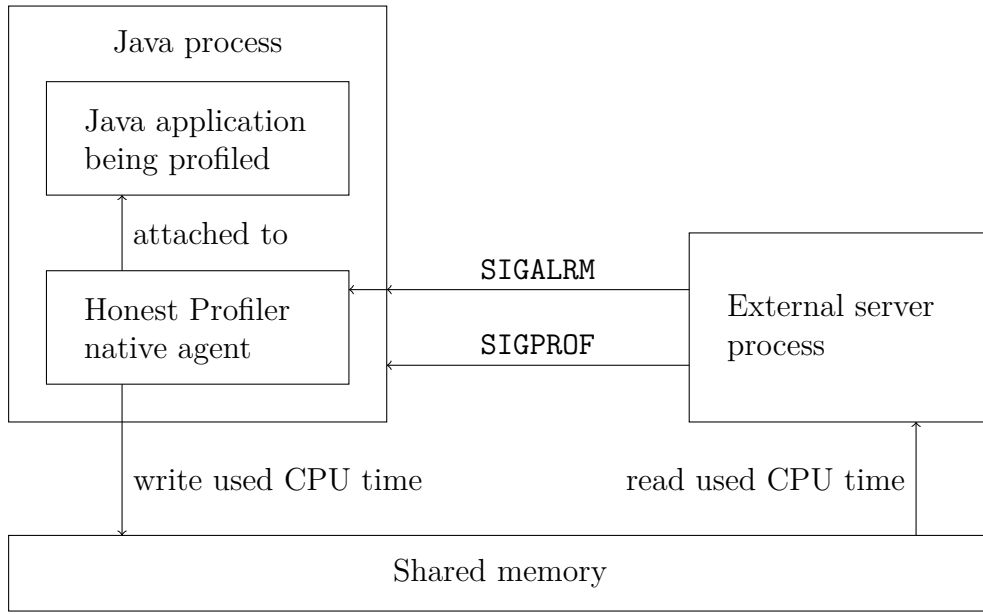


Figure 8: Architecture of the concept

sleeping logic to reduce overhead worth mentioning?

Testing shows that such modification for Honest Profiler runs the benchmark test in 72.64 seconds on average which is better than the baseline benchmark execution time. This can be credited to the fact that time measurement logic has been extracted from the Honest Profiler's logic which means that the Java process will have to execute fewer amount of instructions. From running the benchmark test, a sample was obtained roughly after every 100 microseconds. 746868 samples were obtained during profiling and 38775 (5.2%) of them were useful. This approach obtained 256% more useful samples than the baseline benchmark.

Visualization of the profiling results from the benchmark run using this solution can be observed in figure 9 on page 25.




Figure 9: Visualization of the external timekeeping based solution profiling results

Whereas this approach is slightly less cumbersome to set up than installing a custom kernel as demonstrated in section 4.2.2, the small percentage of useful samples of all samples obtained indicates to problems in this implementation which are briefly covered in the section 5.

5 Future research

This thesis did not focus on evaluating the accuracy of the provided solutions. For future studies, investigating the accuracy of the provided solutions based on the ideas brought out in the Mytkowicz’s paper on profiler accuracy evaluation [11] could improve the practical value of these solutions.

Additionally, during testing the solutions, a problem with large amount of unusable samples occurred. Namely, since the **SIGPROF** signal sent to the profileable process is sent to an *arbitrary* thread then it is possible that the signal is handled by a thread which is not able to provide a call trace sample that could be used in results’ visualization. Unusable sample is obtained when the **AsyncGetCallTrace** result a negative return value which mostly happens in the following cases:

Should present all cases or most common ones would suffice?

1. **SIGPROF** signal is sent to a non-Java thread such as native threads for just-in-time compilation and garbage collection
2. **SIGPROF** signal is sent to a Java thread while garbage collection is active
3. **SIGPROF** signal is sent to a thread which is exiting
4. **SIGPROF** signal is sent to a thread whose stack trace is not walkable

Future studies could investigate the signals emitted and investigate the distribution of threads which are handling the signal and block the threads from handling the profiling signal as they will not provide usable stack traces for profiling results. For example, patching the main thread loop in the Java Virtual Machine initialization using the construct shown in listing 5 decreases the amount of unusable samples by a noticeable margin due to the fact that this thread can not handle the profiling signal sent by Honest Profiler. A promising idea would be to investigate the initialization of native threads within Java Virtual Machine and disable handling of the profiling signal.

```
sigset_t mask;  
sigemptyset(&mask);  
sigaddset(&mask, SIGPROF);  
pthread_sigmask(SIG_BLOCK, &mask, NULL)
```

Listing 5: Blocking the **SIGPROF** signal for a thread

6 Conclusion

The thesis thoroughly explained the ideas behind sampling profiling and described the problems that could affect such profiling method. It then dissects the internals of Honest Profiler and explains why it is a viable choice

and in which use-cases it outshines the others – sampling bias etc

as profiler compared to

other

existing profiling solutions.

First sentence in past, second in present. Also, style question, in the thesis or thesis?

The performance of Honest Profiler's information acquiring logic was measured. Based on the performance measurements, it was apparent that it is possible

there was no fundamental reason why it was not possible?

to increase the amount of information that can be extracted from the observable application by Honest Profiler.

The main result of this thesis are three solutions which increase the profiling information output without causing noticeable performance overhead. The solution's performance and extracted information amount was tested on a benchmark test to verify the results. One of the provided solutions proved to produce reliable results with increased information output and negligible performance overhead. Although two of the provided solutions introduce new problems, ideas for improvement are presented in the future research section 5

Discussed in Slack. I think currently, the two other solutions are undersold.

.

References

- [1] getitimer(2) - linux manual page.
- [2] Honest profiler: A sampling jvm profiler without the safepoint sample bias. <https://github.com/jvm-profiling-tools/honest-profiler>.
- [3] HotSpot glossary of terms.
- [4] jstack - stack trace.
- [5] Jvmtm tool interface. <https://docs.oracle.com/javase/8/docs/platform/jvmti/jvmti.html>.
- [6] JVM(TM) tool interface 1.2.3.
- [7] Openjdk, jdk8: AsyncGetCallTrace source. <http://hg.openjdk.java.net/jdk8/jdk8/hotspot/file/tip/src/share/vm/prims/forte.cpp#l1513>.
- [8] time(7) - linux manual page.
- [9] visualvm 0.20080728, stacks_8c-source.html.
- [10] Brendan Gregg. Flame graphs.
- [11] Todd Mytkowicz, Amer Diwan, Matthias Hauswirth, and Peter F. Sweeney. Evaluating the accuracy of java profilers. 45(6):187.
- [12] W. Richard Stevens and Stephen A. Rago. *Advanced Programming in the UNIX Environment*. Addison-Wesley. Google-Books-ID: kCTMFpEcIOwC.
- [13] Linus Torvalds. linux: Linux kernel source tree. original-date: 2011-09-04T22:48:12Z.
- [14] nitsan wakart. psychosomatic, lobotomy, saw: safepoints: meaning, side effects and overheads.
- [15] Nitsan Wakart. Psychosomatic, lobotomy, saw: Why (most) sampling java profilers are fucking terrible.
- [16] Jack Whitham. Jack whitham.

Appendices

A Source code

All optimization implementations for Honest profiler are available at the following GitHub repository in separate branches:

<https://github.com/Oja95/honest-profiler>

External server implementation for the shared memory solution described in section 4.2.3 is available at the following GitHub repository:

<https://github.com/Oja95/honest-profiler-shared-mem-server>

B Benchmark test code

```
package ee.ut.sorting;

import java.io.IOException;
import java.nio.charset.StandardCharsets;
import java.nio.file.*;
import java.util.stream.Collectors;
import java.util.*;

public class Main {
    public static void main(String[] args) throws IOException {
        long totalStart = System.nanoTime();

        List<Integer> sortableList = Files.readAllLines(Paths.get("/tmp/integers.data"),
            StandardCharsets.UTF_8)
            .stream().mapToInt(x -> Integer.parseInt(x.trim())).boxed().collect(Collectors.toList());

        for (int i = 0; i < 10; i++) {
            QuickSorterImpl<Integer> integerQuickSorter = new QuickSorterImpl<>();
            long start = System.nanoTime();
            integerQuickSorter.sort(new ArrayList<>(sortableList));
            long end = System.nanoTime() - start;
            System.out.println(end);
        }
        long totalEnd = System.nanoTime() - totalStart;
        System.out.println(totalEnd);
    }
}
```

Non-exclusive licence to reproduce thesis and make thesis public

I, **Tiit Oja**,

1. herewith grant the University of Tartu a free permit (non-exclusive licence) to:
 - 1.1. reproduce, for the purpose of preservation and making available to the public, including for addition to the DSpace digital archives until expiry of the term of validity of the copyright, and
 - 1.2. make available to the public via the web environment of the University of Tartu, including via the DSpace digital archives until expiry of the term of validity of the copyright,

Optimizing JVM profiling performance for Honest Profiler,

supervised by **Vootele Rõtov** and **Vesal Vojdani**,

2. I am aware of the fact that the author retains these rights.
3. I certify that granting the non-exclusive licence does not infringe the intellectual property rights or rights arising from the Personal Data Protection Act.

Tartu, dd.mm.yyyy

Add date